

130050027 (Anmol Arora), 130050041 (Aman Goel)

Lab 8 report

Part 1:

The maximum throughput observed is **5.879** files/sec, which is observed for the first time at **N=2**. Thus the number of worker threads at the server to fully saturate the server is **2** {The difference in throughput at N=1 and N=2 is insignificant, owing due to small calculation errors, hence for practical purposes the saturation is reached at N=1 itself}.

The load created by 20 users is small enough to be handled by a couple of server threads, hence the saturation is reached pretty quickly. The bottleneck resource here is the network bandwidth as apparent on running iftop, top and iostat:

```
top - 11:24:27 up 3 days, 19:23,  3 users,  load average: 0.80, 0.61, 0.50
Tasks: 256 total,  2 running, 254 sleeping,  0 stopped,  0 zombie
%Cpu(s): 13.6 us,  0.8 sy,  0.0 ni, 83.1 id,  1.4 wa,  0.0 hi,  1.1 si,  0.0 st
KiB Mem:  8054896 total, 7868552 used, 186344 free, 119396 buffers
KiB Swap:  0 total,  0 used,  0 free. 3204744 cached Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
30163	root	20	0	457288	1172	964	S	59.5	0.0	0:56.29	server-mt
2406	mewtomat	20	0	2276544	443064	73140	S	0.7	5.5	61:18.53	chrome
8	root	20	0	0	0	0	S	0.3	0.0	0:45.15	rcuos/0
1396	root	20	0	523180	176576	160940	S	0.3	2.2	39:39.32	Xorg

This shows CPU is not the bottleneck.

```
mewtomat@mewtomat-Inspiron-5737:~$ iostat -d 1 -m
Linux 3.13.0-68-generic (mewtomat-Inspiron-5737)           Thursday 10 March 2016
```

Device:	tps	MB_read/s	MB_wrtn/s	MB_read	MB_wrtn
sda	3.75	0.03	0.05	10433	16206
Device:	tps	MB_read/s	MB_wrtn/s	MB_read	MB_wrtn
sda	0.00	0.00	0.00	0	0
Device:	tps	MB_read/s	MB_wrtn/s	MB_read	MB_wrtn
sda	0.00	0.00	0.00	0	0
Device:	tps	MB_read/s	MB_wrtn/s	MB_read	MB_wrtn
sda	0.00	0.00	0.00	0	0
Device:	tps	MB_read/s	MB_wrtn/s	MB_read	MB_wrtn
sda	0.00	0.00	0.00	0	0

This shows disk is not the bottleneck.

	203Mb	406Mb	610Mb	813Mb	0.99Gb
192.168.0.104		=> 192.168.0.100		94.7Mb	93.3Mb
		<=		1.77Mb	1.75Mb
192.168.0.104		=> 192.168.0.1		756b	1.28kb
		<=		712b	938b
192.168.0.104		=> 192.168.0.255		0b	333b
		<=		0b	0b
192.168.0.104		=> 10.5.100.250		0b	262b
		<=		0b	0b
192.168.0.104		=> 255.255.255.255		0b	166b
		<=		0b	0b
192.168.0.101		=> 192.168.0.255		0b	0b
		<=		0b	0b
192.168.0.255		=> 192.168.0.100		0b	0b
		<=		0b	0b
255.255.255.255		=> 192.168.0.100		0b	0b
		<=		0b	0b
<div> <div></div> <div>TX: cum: 854MB peak: 94.7Mb</div> <div>RX: 16.9MB 1.93Mb</div> <div>TOTAL: 871MB 96.5Mb</div> </div> <div> <div></div> <div>rates: 94.7Mb 93.3Mb 93.0Mb</div> <div>96.5Mb 95.0Mb 94.8Mb</div> </div>					

The throughput measurements for part 1 are as follows:

Part 2:

When the experiment is run for large number of users, it is observed that many clients are denied service. This is reflected by the error message “Error: Connection Timed Out” being printed multiple times at the client. This is happening because the client program is sending a large number of requests at a rate which is greater than the rate of servicing at the server. As a result the request queue at server gets full and it drops any more requests from the client. The clients corresponding to these dropped requests wait for some time, hoping to get a reply from server, but eventually get timed out and print the above mentioned error. These are the client threads which are being denied service from server.

On measuring the throughput values for each N from 1 to 10, it is observed that initially the throughput increases but then it decreases as N is increased further.

The increase in throughput can be explained from the fact that with greater number of worker threads, the requests in queue get served quickly, enabling more and more pending requests to get accepted and subsequently getting served, giving higher throughput.

The decrease in throughput for further higher values of N is because of contention among threads for the lock. The threads acquire lock at various points of time in their runtime. At a time only one of the threads can hold a lock. But when there are many threads waiting for the same lock, the whole program starts incurring significant amounts of overheads in lock acquiring. Thus the latency of serving the requests starts increasing at a higher rate than the gain in throughput through increased successful requests, resulting in overall throughput decrease for large N.

Measurements for part 2 are as follows:

n = 1

Total Successful Requests: 620

Throughput: 2.846

n = 2

Total Successful Requests: 691

Throughput: 3.509

n = 3

Total Successful Requests: 747

Throughput: 4.086

n = 4

Total Successful Requests: 816

Throughput: 4.282

n = 5

Total Successful Requests: 812

Throughput: 4.364

n = 6

Total Successful Requests: 812

Throughput: 4.533

n = 7

Total Successful Requests: 803

Throughput: 4.435

n = 8

Total Successful Requests: 838

Throughput: 4.191

n = 9

Total Successful Requests: 811

Throughput: 4.004

n = 10

Total Successful Requests: 836

Throughput: 3.674

Plot of Throughput vs N for Part 2

